

telco_churn_model

February 2, 2023

1 Final Project Submission

Student name: David Boyd

Student pace: Self paced

Scheduled project review date/time: 3/02/2023 19:00 UK time

Instructor name: Abhineet

1.1 Business Objective

A telecommunications company wants to identify the leading factors of why a customer cancels their service. By understanding the factors which most strongly contribute to their churn rate they can set up work streams to help mitigate and reduce this effect over time, which will increase the LTV and overall ROI for the company.

1.2 Importing packages & EDA

```
[44]: # Packages that I'm using to create the classification
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV,
↳StratifiedKFold
from sklearn.metrics import classification_report, precision_score,
↳recall_score, accuracy_score, f1_score
from sklearn.metrics import plot_confusion_matrix, confusion_matrix,
↳plot_roc_curve
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier,
↳GradientBoostingClassifier
from sklearn.preprocessing import StandardScaler, OneHotEncoder,
↳FunctionTransformer
from sklearn.compose import ColumnTransformer
```

```

from sklearn.pipeline import Pipeline, FeatureUnion
from sklearn.linear_model import LogisticRegression

from xgboost import XGBClassifier

from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline as ImPipeline

import warnings
warnings.filterwarnings('ignore')

```

```

[45]: # Loads data into a DataFrame for investigation
df = pd.read_csv('data/telco_churn_data.csv')

df.head()

```

```

[45]:  state  account length  area code phone number international plan \
0     KS             128      415    382-4657                no
1     OH             107      415    371-7191                no
2     NJ             137      415    358-1921                no
3     OH              84      408    375-9999                yes
4     OK              75      415    330-6626                yes

      voice mail plan  number vmail messages  total day minutes  total day calls \
0                yes                25          265.1          110
1                yes                26          161.6          123
2                no                 0          243.4          114
3                no                 0          299.4           71
4                no                 0          166.7          113

      total day charge  ...  total eve calls  total eve charge \
0          45.07  ...           99          16.78
1          27.47  ...          103          16.62
2          41.38  ...          110          10.30
3          50.90  ...           88           5.26
4          28.34  ...          122          12.61

      total night minutes  total night calls  total night charge \
0          244.7           91          11.01
1          254.4          103          11.45
2          162.6          104           7.32
3          196.9           89           8.86
4          186.9          121           8.41

      total intl minutes  total intl calls  total intl charge \
0           10.0           3           2.70

```

1	13.7	3	3.70
2	12.2	5	3.29
3	6.6	7	1.78
4	10.1	3	2.73

	customer service calls	churn
0	1	False
1	1	False
2	0	False
3	2	False
4	3	False

[5 rows x 21 columns]

```
[46]: # Checking what datatypes each column consists of
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   state                                3333 non-null   object
1   account length                       3333 non-null   int64
2   area code                           3333 non-null   int64
3   phone number                        3333 non-null   object
4   international plan                  3333 non-null   object
5   voice mail plan                     3333 non-null   object
6   number vmail messages               3333 non-null   int64
7   total day minutes                   3333 non-null   float64
8   total day calls                     3333 non-null   int64
9   total day charge                     3333 non-null   float64
10  total eve minutes                   3333 non-null   float64
11  total eve calls                     3333 non-null   int64
12  total eve charge                     3333 non-null   float64
13  total night minutes                 3333 non-null   float64
14  total night calls                   3333 non-null   int64
15  total night charge                  3333 non-null   float64
16  total intl minutes                  3333 non-null   float64
17  total intl calls                    3333 non-null   int64
18  total intl charge                   3333 non-null   float64
19  customer service calls              3333 non-null   int64
20  churn                              3333 non-null   bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB
```

```
[47]: # Converting target variable into an integer for ease of use later
df['churn'] = df['churn'].astype('int')
```

```
[48]: #
df.describe()
```

```
[48]:
```

	account length	area code	number vmail messages	total day minutes	\
count	3333.000000	3333.000000	3333.000000	3333.000000	
mean	101.064806	437.182418	8.099010	179.775098	
std	39.822106	42.371290	13.688365	54.467389	
min	1.000000	408.000000	0.000000	0.000000	
25%	74.000000	408.000000	0.000000	143.700000	
50%	101.000000	415.000000	0.000000	179.400000	
75%	127.000000	510.000000	20.000000	216.400000	
max	243.000000	510.000000	51.000000	350.800000	

	total day calls	total day charge	total eve minutes	total eve calls	\
count	3333.000000	3333.000000	3333.000000	3333.000000	
mean	100.435644	30.562307	200.980348	100.114311	
std	20.069084	9.259435	50.713844	19.922625	
min	0.000000	0.000000	0.000000	0.000000	
25%	87.000000	24.430000	166.600000	87.000000	
50%	101.000000	30.500000	201.400000	100.000000	
75%	114.000000	36.790000	235.300000	114.000000	
max	165.000000	59.640000	363.700000	170.000000	

	total eve charge	total night minutes	total night calls	\
count	3333.000000	3333.000000	3333.000000	
mean	17.083540	200.872037	100.107711	
std	4.310668	50.573847	19.568609	
min	0.000000	23.200000	33.000000	
25%	14.160000	167.000000	87.000000	
50%	17.120000	201.200000	100.000000	
75%	20.000000	235.300000	113.000000	
max	30.910000	395.000000	175.000000	

	total night charge	total intl minutes	total intl calls	\
count	3333.000000	3333.000000	3333.000000	
mean	9.039325	10.237294	4.479448	
std	2.275873	2.791840	2.461214	
min	1.040000	0.000000	0.000000	
25%	7.520000	8.500000	3.000000	
50%	9.050000	10.300000	4.000000	
75%	10.590000	12.100000	6.000000	
max	17.770000	20.000000	20.000000	

	total intl charge	customer service calls	churn
--	-------------------	------------------------	-------

count	3333.000000	3333.000000	3333.000000
mean	2.764581	1.562856	0.144914
std	0.753773	1.315491	0.352067
min	0.000000	0.000000	0.000000
25%	2.300000	1.000000	0.000000
50%	2.780000	1.000000	0.000000
75%	3.270000	2.000000	0.000000
max	5.400000	9.000000	1.000000

```
[49]: df.isna().sum()
```

```
[49]: state                0
account length           0
area code                0
phone number            0
international plan       0
voice mail plan          0
number vmail messages    0
total day minutes        0
total day calls           0
total day charge          0
total eve minutes        0
total eve calls           0
total eve charge          0
total night minutes      0
total night calls         0
total night charge        0
total intl minutes       0
total intl calls          0
total intl charge         0
customer service calls    0
churn                    0
dtype: int64
```

```
[50]: # Creating area code into a categorical variable, instead of int as it's more
      ↳ applicable
df['area code'] = df['area code'].astype('str')
```

```
[51]: df.churn.value_counts()
```

```
[51]: 0    2850
      1     483
      Name: churn, dtype: int64
```

Looking at the above we can see that this dataset is imbalanced, therefore in the process to maximise the model performance on unseen data we'll have to utilise a resampling method on the training set.

```
[52]: # observing how many and what different options are available for the
      ↪categorical features
df_obj_cols = df.select_dtypes(include='object')

for col in df_obj_cols:
    print(f'{col.title()} has the following different values: {df[col].
    ↪unique()}')
    print(f'There are {df[col].nunique()} options for the {col.title()} feature')
```

```
State has the following different values: ['KS' 'OH' 'NJ' 'OK' 'AL' 'MA' 'MO'
'LA' 'WV' 'IN' 'RI' 'IA' 'MT' 'NY'
'ID' 'VT' 'VA' 'TX' 'FL' 'CO' 'AZ' 'SC' 'NE' 'WY' 'HI' 'IL' 'NH' 'GA'
'AK' 'MD' 'AR' 'WI' 'OR' 'MI' 'DE' 'UT' 'CA' 'MN' 'SD' 'NC' 'WA' 'NM'
'NV' 'DC' 'KY' 'ME' 'MS' 'TN' 'PA' 'CT' 'ND']
```

There are 51 options for the State feature

```
Area Code has the following different values: ['415' '408' '510']
```

There are 3 options for the Area Code feature

```
Phone Number has the following different values: ['382-4657' '371-7191'
'358-1921' ... '328-8230' '364-6381' '400-4344']
```

There are 3333 options for the Phone Number feature

```
International Plan has the following different values: ['no' 'yes']
```

There are 2 options for the International Plan feature

```
Voice Mail Plan has the following different values: ['yes' 'no']
```

There are 2 options for the Voice Mail Plan feature

Looking at the above columns, it is clear that while we have been give a phone number, since these are all unique, this isn't going to have any influence over whether a customer has churned or not, therefore the next steps will be to drop this column.

After this, looking at our brief exploratory data above there are a few different steps to complete before the data is ready to be passed into a model. These are: - One hot encode all of the categorical columns - Normalize all of the numerical columns - Account for the class imbalance within the target data by using SMOTE

Once these steps have been completed the data should be ready to pass into our proposed models

1.3 Setting up datasets for model testing

```
[53]: # Removing phone number as it'll lead to overfitting given each number is unique
X = df.drop(columns=['churn', 'phone number'], axis=1)

# Selecting target variable
y = df.churn
```

```
[54]: # Splitting the data for model testing
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
↳stratify=y, random_state=42)
```

1.4 Useful Functions

Below I've created a few functions that are useful when it comes to building out the model. The first is focused on adding additional features through feature engineering to enrich the dataset

```
[55]: def additional_churn_features(X):

    X['total_calls'] = X['total day calls'] + X['total eve calls'] + X['total_
↳night calls'] + X['total intl calls']
    X['total_charges'] = X['total day charge'] + X['total eve charge'] +
↳X['total night charge'] + X['total intl charge']
    X['total_minutes'] = X['total day minutes'] + X['total eve minutes'] +
↳X['total night minutes'] + X['total intl minutes']
    X['pct_intl_calls'] = X['total intl calls'] / X['total_calls']
    X['pct_domestic_calls'] = 1 - X['pct_intl_calls']
    X['pct_intl_minutes'] = X['total intl minutes'] / X['total_minutes']
    X['pct_intl_charges'] = X['total intl charge'] / X['total_charges']
    X['avg_mins_per_call_day'] = X['total day minutes'] / X['total day calls']
    X['avg_mins_per_call_eve'] = X['total eve minutes'] / X['total eve calls']
    X['avg_mins_per_call_night'] = X['total night minutes'] / X['total night_
↳calls']

    X.fillna(method='bfill', inplace=True)

    return X
```

```
[56]: def preparing_model_pipeline(X, model=DecisionTreeClassifier, grid=None,
↳score='f1', booster=None):
    """
    This function takes a dataset, a model to instantiate if provided and a grid
    which then creates a sklearn pipeline to perform the following steps:
    - Adding additional features to the dataset
    - OHE all categorical features
    - Scaling all numeric columns
    - performing SMOTE to account for class imbalance
    - Instantiating the model

    If grid is provided, it then performs the relevant gridsearch function
    """
    # feature engineering
    feat_eng_transformer = FunctionTransformer(additional_churn_features,
↳validate=False)
```

```

# OHE
cat_features = X.select_dtypes(exclude=[np.number]).columns.tolist()
cat_col_transformer = ColumnTransformer(transformers=[
    ("ohe", OneHotEncoder(handle_unknown='ignore', sparse=True),
    ↪cat_features)], remainder="passthrough")

# SMOTE & StandardScaler
if booster == None:
    imb_pipe = ImPipeline(steps=[('feat_eng', feat_eng_transformer),
                                ('cat_ohe', cat_col_transformer),
                                ('sscaler', StandardScaler(with_mean=False)),
                                ('smote', SMOTE(random_state=42)),
                                ('model', model(random_state=42))
                                ])
else:
    imb_pipe = ImPipeline(steps=[('feat_eng', feat_eng_transformer),
                                ('cat_ohe', cat_col_transformer),
                                ('sscaler', StandardScaler(with_mean=False)),
                                ('smote', SMOTE(random_state=42)),
                                ('model', model(random_state=42, booster=booster))
                                ])

if grid == None:
    return imb_pipe

else:
    stratified_kfold = StratifiedKFold(n_splits=3,
                                       shuffle=True,
                                       random_state=42)

    gridsearch = GridSearchCV(estimator=imb_pipe,
                              param_grid=grid,
                              scoring=score,
                              cv=stratified_kfold)

    return gridsearch

```

1.5 Set up baseline performance for all models

In the code below we create a dataframe to show what performance looks like across a range of different classifier models at a baseline level, this will allow us to select a few to look at hypertuning using the GridSearchCV method


```
[57]: model_list = [DecisionTreeClassifier, RandomForestClassifier,
↳ LogisticRegression,
        AdaBoostClassifier, GradientBoostingClassifier, XGBClassifier]
model_performance = []
for model in model_list:
    inner_list = []
    inner_list.append(f'{str(model).split(".")[1][:-2]}')

    result = preparing_model_pipeline(X_train, model)

    fitted_model = result.fit(X_train, y_train)

    recall_train = recall_score(y_train, fitted_model.predict(X_train))
    recall_test = recall_score(y_test, fitted_model.predict(X_test))
    inner_list.append(recall_train)
    inner_list.append(recall_test)

    f1_train = f1_score(y_train, fitted_model.predict(X_train))
    f1_test = f1_score(y_test, fitted_model.predict(X_test))
    inner_list.append(f1_train)
    inner_list.append(f1_test)

    model_performance.append(inner_list)

    #print(f'This is a {model} model')
    #print(classification_report(y_test, fitted_model.predict(X_test), digits=3))

base_model_df = pd.DataFrame(model_performance, columns= ['model_name',
↳ 'recall_train', 'recall_test', 'f1_train', 'f1_test'])

base_model_df
```

```
[57]:
```

	model_name	recall_train	recall_test	f1_train	f1_test
0	DecisionTreeClassifier	1.000000	0.820690	1.000000	0.725610
1	RandomForestClassifier	1.000000	0.744828	1.000000	0.791209
2	LogisticRegression	0.766272	0.662069	0.516451	0.446512
3	AdaBoostClassifier	0.730769	0.648276	0.704708	0.620462
4	GradientBoostingClassifier	0.875740	0.813793	0.933754	0.877323
5	XGBClassifier	1.000000	0.834483	1.000000	0.886447

Looking at the above, we'll select the following models to tune further: - DecisionTreeClassifier - RandomForestClassifier - XGBClassifier

These will run through a grid search and return both the best parameters for performance metrics, targetting the f1 metric

```
[15]: grid_dt = [{'model__criterion': ['gini', 'entropy'],
    'model__max_depth': [None, 5, 10, 15, 20],
    'model__max_features': [None, 10, 5, 15],
    'model__min_samples_split': [2, 5, 10],
    'model__min_samples_leaf': [1, 2, 5]}]

tuned_decision_model = preparing_model_pipeline(X_train, grid=grid_dt)

fitted_tuned_dt = tuned_decision_model.fit(X_train, y_train)

fitted_tuned_dt.best_params_
```

```
[15]: {'model__criterion': 'gini',
    'model__max_depth': 5,
    'model__max_features': None,
    'model__min_samples_leaf': 1,
    'model__min_samples_split': 2}
```

```
[16]: pd.DataFrame(classification_report(y_test, fitted_tuned_dt.predict(X_test),
    ↳digits=3, output_dict=True)).transpose()
```

```
[16]:
```

	precision	recall	f1-score	support
0	0.953462	0.982456	0.967742	855.000
1	0.873950	0.717241	0.787879	145.000
accuracy	0.944000	0.944000	0.944000	0.944
macro avg	0.913706	0.849849	0.877810	1000.000
weighted avg	0.941933	0.944000	0.941662	1000.000

```
[17]: grid_rf = {
    'model__n_estimators': [100, 150],
    'model__criterion': ['entropy', 'gini'],
    'model__max_depth': [3, 7],
    'model__max_features': [5, 10, 15],
    'model__min_samples_split': [10, 20, 50],
    'model__min_samples_leaf': [1, 2, 4]
}

tuned_rforest_model = preparing_model_pipeline(X_train,
    ↳model=RandomForestClassifier, grid=grid_rf)

fitted_tuned_rf = tuned_rforest_model.fit(X_train, y_train)

fitted_tuned_rf.best_params_
```

```
[17]: {'model__criterion': 'entropy',
    'model__max_depth': 7,
    'model__max_features': 15,
    'model__min_samples_leaf': 4,
```

```
'model__min_samples_split': 10,
'model__n_estimators': 100}
```

```
[18]: pd.DataFrame(classification_report(y_test, fitted_tuned_rf.predict(X_test),
↳digits=3, output_dict=True)).transpose()
```

```
[18]:
```

	precision	recall	f1-score	support
0	0.968531	0.971930	0.970228	855.000
1	0.830986	0.813793	0.822300	145.000
accuracy	0.949000	0.949000	0.949000	0.949
macro avg	0.899759	0.892861	0.896264	1000.000
weighted avg	0.948587	0.949000	0.948778	1000.000

```
[58]: grid_xgb = {
    'model__learning_rate': [0.1, 0.2],
    'model__max_depth': [6],
    'model__min_child_weight': [1, 2],
    'model__subsample': [0.5, 0.7],
    'model__n_estimators': [100],
}

tuned_xgboost_model = preparing_model_pipeline(X_train, model=XGBClassifier,
↳grid=grid_xgb, score='recall', booster='gbtree')

tuned_xgboost_model.fit(X_train, y_train)

tuned_xgboost_model.best_params_
```

```
[58]: {'model__learning_rate': 0.1,
'model__max_depth': 6,
'model__min_child_weight': 1,
'model__n_estimators': 100,
'model__subsample': 0.7}
```

```
[59]: pd.DataFrame(classification_report(y_test, tuned_xgboost_model.predict(X_test),
↳digits=3, output_dict=True)).transpose()
```

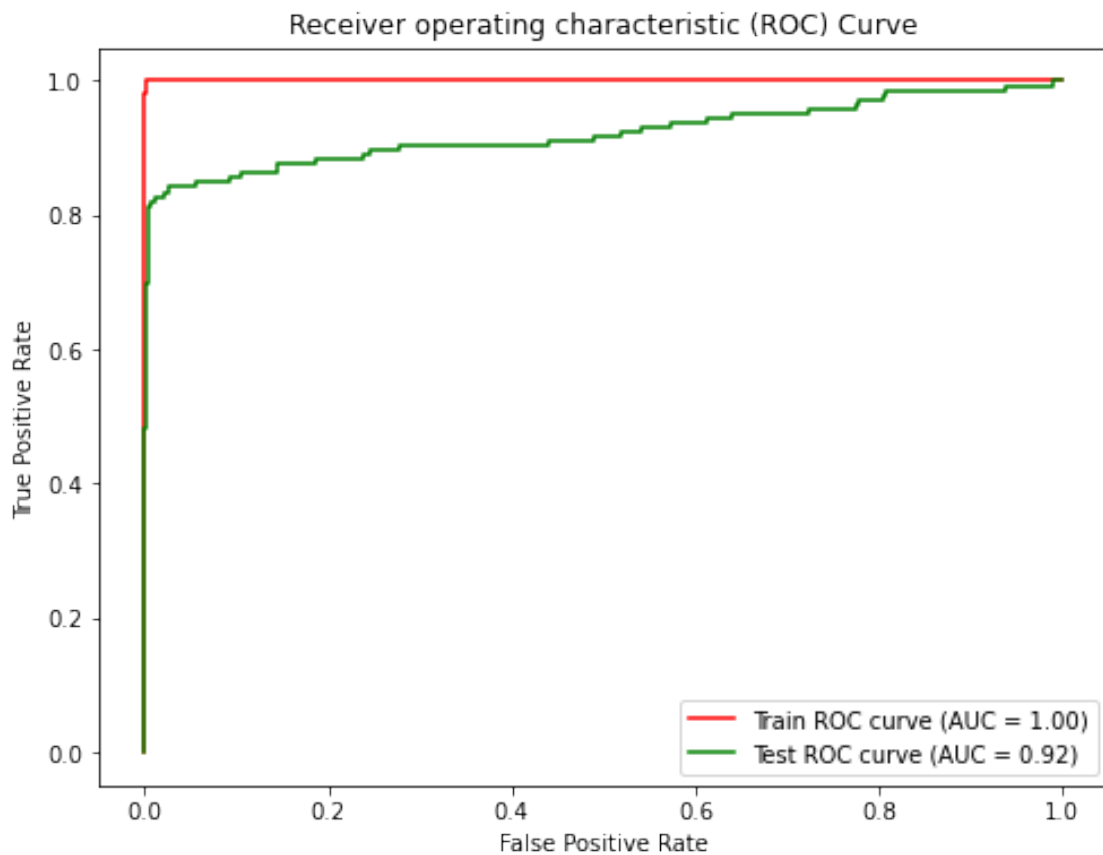
```
[59]:
```

	precision	recall	f1-score	support
0	0.969213	0.994152	0.981524	855.000
1	0.959350	0.813793	0.880597	145.000
accuracy	0.968000	0.968000	0.968000	0.968
macro avg	0.964281	0.903973	0.931061	1000.000
weighted avg	0.967783	0.968000	0.966890	1000.000

Now we've performed the tuning on each model, it appears that the best tradeoff between recall and f1-score is with our XGBoostClassifier model. As a result, below we'll produce a few different plots.


The first will be plotting our AUC-ROC curve After that we want to plot the confusion matrix on our test dataset

```
[60]: fig, ax2 = plt.subplots(figsize=(8,6))
plot_roc_curve(tuned_xgboost_model, X_train, y_train, ax=ax2, name = 'Train ROC_
↪curve', color='r')
plot_roc_curve(tuned_xgboost_model, X_test, y_test, ax=ax2, name = 'Test ROC_
↪curve', color='g')
ax2.set_xlabel('False Positive Rate')
ax2.set_ylabel('True Positive Rate')
ax2.set_title('Receiver operating characteristic (ROC) Curve')
plt.show();
```



```
[61]: fig, ax = plt.subplots(figsize=(8,5))

plot_confusion_matrix(tuned_xgboost_model, X_test, y_test, ax=ax,
↪normalize='true')
ax.set_title("XGBoost Confusion Matrix - Test");
```




Below we look to extract the name of the features which have the highest importance in the selected model, to highlight these back to our stakeholders for areas that need further investigation

```
[62]: # Identifying which features have the highest level of importance

feat_import_array = tuned_xgboost_model.estimator.fit(X_train, y_train).
    ↳ steps[4][1].feature_importances_

[67]: imp_features = pd.DataFrame(pd.Series(feat_import_array, index=X_train_df.
    ↳ columns).sort_values(ascending=False))
imp_features = imp_features.head(10)
imp_features['index'] = imp_features.index
imp_features.head()

ax = sns.barplot(x=imp_features[0], y=imp_features['index'], data=imp_features)
ax.set_xlabel('Importance')
ax.set_ylabel('Feature')
ax.set_title('Importance per Feature');
```



1.6 Section to be removed (Code used to preprocess without a pipeline)

The below code is used in the section above to extract the correct column names for feature importance. As we've combined our model into the pipeline, it doesn't currently extract the DataFrame with the preprocessed features for more efficient code.

```
[63]: X_train_cat = X_train.select_dtypes(include='object')

ohe = OneHotEncoder(handle_unknown='ignore', sparse=False)

X_train_cat_encoded = ohe.fit_transform(X_train_cat)
X_train_cat_encoded_df = pd.DataFrame(X_train_cat_encoded,
                                      columns=ohe.get_feature_names(),
                                      index=X_train_cat.index)
X_train_cat_encoded_df.shape
```

[63]: (2333, 58)

```
[64]: X_train_cat_encoded_df
```

```
[64]:
```

	x0_AK	x0_AL	x0_AR	x0_AZ	x0_CA	x0_CO	x0_CT	x0_DC	x0_DE	x0_FL	\
606	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
2468	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
1844	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
3187	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

3083	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
...
2670	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2165	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2988	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
179	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2762	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0

	...	x0_WI	x0_WV	x0_WY	x1_408	x1_415	x1_510	x2_no	x2_yes	x3_no	\
606	...	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	1.0	
2468	...	0.0	1.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	
1844	...	0.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	0.0	
3187	...	0.0	1.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	
3083	...	0.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	1.0	
...	
2670	...	0.0	0.0	1.0	0.0	0.0	1.0	1.0	0.0	0.0	
2165	...	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	1.0	
2988	...	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	1.0	
179	...	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	
2762	...	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	

	x3_yes
606	0.0
2468	1.0
1844	1.0
3187	1.0
3083	0.0
...	...
2670	1.0
2165	0.0
2988	0.0
179	0.0
2762	0.0

[2333 rows x 58 columns]

```
[65]: X_train_clean = X_train.drop(columns=X_train_cat.columns.values, axis=1)
X_train_df = pd.concat([X_train_cat_encoded_df, X_train_clean], axis=1)

X_train_df
```

[65]:	x0_AK	x0_AL	x0_AR	x0_AZ	x0_CA	x0_CO	x0_CT	x0_DC	x0_DE	x0_FL	\
606	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
2468	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
1844	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
3187	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
3083	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

...
2670	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2165	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2988	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
179	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2762	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0

	...	total_calls	total_charges	total_minutes	pct_intl_calls	\
606	...	329	57.67	571.3	0.024316	
2468	...	306	60.01	650.3	0.006536	
1844	...	296	56.33	599.4	0.010135	
3187	...	259	70.27	702.2	0.015444	
3083	...	387	69.77	707.0	0.012920	

...
2670	...	312	60.90	560.0	0.019231	
2165	...	319	61.33	605.9	0.009404	
2988	...	234	72.22	668.9	0.021368	
179	...	346	73.36	725.6	0.000000	
2762	...	272	32.89	394.4	0.014706	

	pct_domestic_calls	pct_intl_minutes	pct_intl_charges	\
606	0.975684	0.022930	0.061384	
2468	0.993464	0.006612	0.019330	
1844	0.989865	0.019186	0.055210	
3187	0.984556	0.007690	0.020777	
3083	0.987080	0.017680	0.048445	

...
2670	0.980769	0.016071	0.039901
2165	0.990596	0.021456	0.057231
2988	0.978632	0.017940	0.044863
179	1.000000	0.000000	0.000000
2762	0.985294	0.024087	0.078139

	avg_mins_per_call_day	avg_mins_per_call_eve	avg_mins_per_call_night
606	1.172269	2.755238	1.334021
2468	1.146875	2.975000	2.670000
1844	1.161538	2.397059	2.804054
3187	5.297500	3.610526	1.514388
3083	2.180556	0.995105	2.417557
...
2670	2.046296	1.279661	2.237500
2165	1.662136	2.012397	1.936957
2988	2.701042	1.806186	6.177778
179	1.902459	2.609821	1.796429
2762	0.572222	1.673469	2.117500

[2333 rows x 83 columns]

1.7 Summary & Main results

From the feature importance chart above, we can see the areas that have the biggest influence on customer churn are: - Not having an international plan in their account - Not having a voice mail plan in their account - The total charges per account - The number of contacts being made to customer services

To set up a clear list of actions for the company, it is now worth performing a more thorough and in depth EDA around those features to identify what thresholds the business wants to set for their operations teams to try and achieve with the goal of reducing overall customer churn.

After these features, we can also see from the dataset that specific area codes and states have a higher influence than others. While it might be worth investigating this aspect, it is worth having caution what actions can be made around this discovery.